

Cloud Computing Assignment 2



Container Practice

2017/04/14
1 / 18



Outline

- OS-level Virtualization
- Namespaces
- Runc: A lightweight universal container runtime
- Assignment Requirements



A Survey on Virtualization Technologies

Susanta Nanda Tzi-cker Chiueh
{susanta,chiueh}@cs.sunysb.edu
Department of Computer Science
SUNY at Stony Brook
Stony Brook, NY 11794-4400



Quote 1


Virtuality differs from **reality** only in the **formal** world, while possessing a similar essence or effect. In the computer world, a virtual environment is perceived the **same** as that of a real environment by application programs and the rest of the world, though the underlying mechanisms are **formally different**.



Quote 2

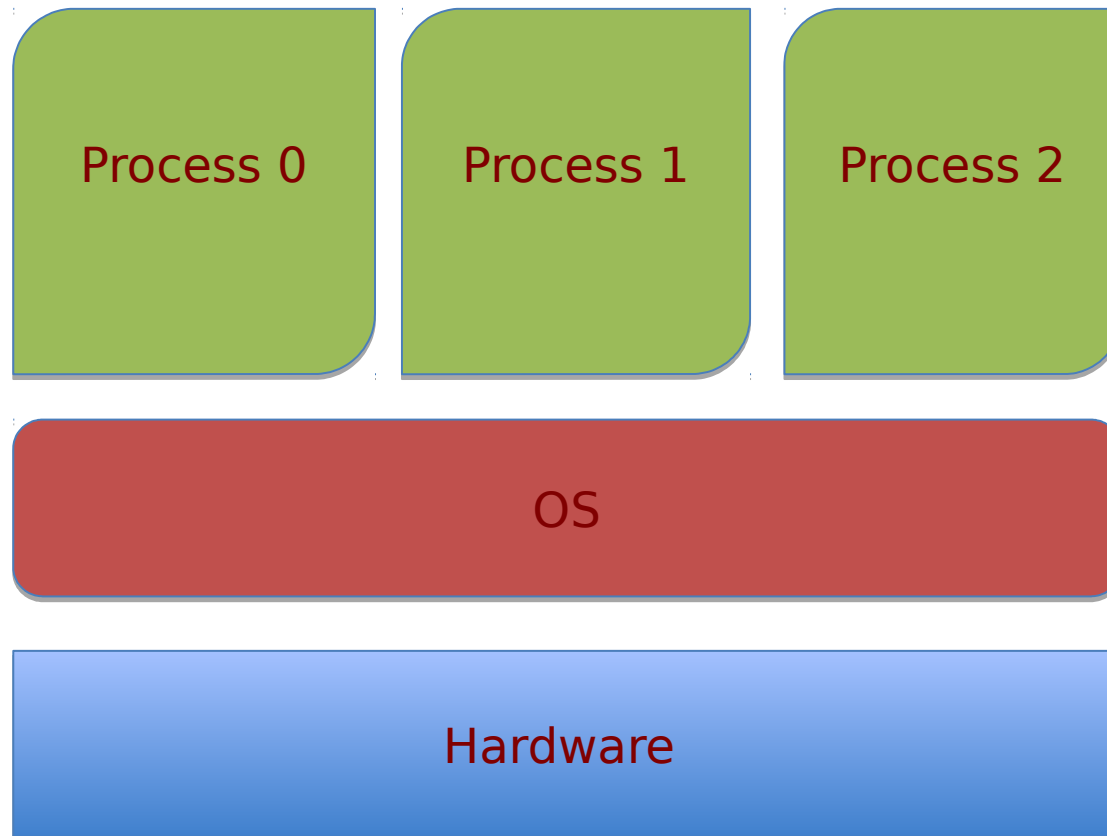
Virtualization is a technology that **combines** or **divides** computing **resources** to present one or many operating environments using methodologies like hardware and software partitioning or aggregation, partial or complete machine simulation, emulation, time-sharing, and many others.

Virtualization (the one you know)




Virtualization?

(the one you familiar with)



OS-level Virtualization (Containerization)



OS-level Virtualization

- In Unix: chroot
- In *BSD: jails
 - why do you “JB” your iphone?
- In Linux: container
 - Namespaces
 - Cgroups
 - Layered File System (optional)



Namespaces

- UTS: Unix Time-sharing System
- PID: Process Identifier
- MNT: Mount
- NET: Networking
- IPC: Inter-Process Communication
- USER
- CGROUP



RunC: The Tutorial

- <https://katacoda.com/lalyos/scenarios/runc>
- We will provide you
 - the **rootfs directory** (in step 2), and
 - the **config.json** (in step 3)for this assignment



Objective: Bridging Two Containers

- Container A and B
- A **simple client** runs on container A, in which user can types arbitrary messages
- An **echo server** runs on container B
- A **bridge** program
- Three programs in this assignment

The Bridge Program

- Shares the **X** namespace with container A, and
- Shares the **Y** namespace with container B
- X and Y are **different ones** of the three namespaces: NET, MNT, IPC
 - network interfaces of different subdomain
 - different root directory
- You may choose **any** of the combinations of X,Y



Environment

- You may use the VM you created in the assignment 1, or
- You may setup a Linux VM on your own machine, and setup runc on it
- For Ubuntu users, **apt-get install runc** should work



Scoring

- Connection between the programs (60%)
 - X namespace connection (20%)
 - Y namespace connection (20%)
 - The echo service work (20%)
- Report (40%)
- Deadline: 5/19 (flexible)



Appendix -- network connection

- Socket
- Bind
- Listen
- Connect
- Accept
- Send, Recv



Appendix -- IPC (msg only) connection

- Msgget
- Msgsnd
- Msgrcv
- Msgctl



Appendix -- file system monitoring

- `inotify_init`
- `inotify_add_watch`